

ArbiSim Guard

Pre-flight transaction simulation for AI agents - working beta

The problem

Injective has built the execution side for agents: the MCP Server, iAgent, the ERC-8004 registry, iBuild. What does not exist yet is the check before execution. An agent can fire a trade that reverts, gets sandwiched, or drains an allowance, and nothing stops it. Autonomous capital, no pre-flight check.

What ArbiSim does

The agent calls ArbiSim before sending a transaction. ArbiSim forks live chain state into an ephemeral Anvil instance, runs the transaction inside the fork, and returns APPROVED or REJECTED with flags: reverts, slippage, MEV/sandwich exposure, gas exhaustion, unsafe allowances, oracle failures. One extra tool call, under a cent (x402, \$0.001 USDC per simulation), and nothing hits the chain until it passes.

Honest status: working beta, not finished

- Live: fork engine on Avalanche Fuji and Arbitrum Sepolia, SimulationRegistry on-chain (Snowtrace verified; Arbitrum Sepolia 0x5Dfd08c3d44BEBfa61a24Af8c2EfbDB5A01dFA32), ERC-4337 UserOp simulation, ERC-8004 reputation checks inside forks, x402 billing, MCP tool preflight_simulate (Claude, Cursor, Eliza, LangGraph).
- In progress: agent deployment gates are still failing in some scenarios and I am fixing them one by one. The UI is overloaded and needs an overhaul. No mainnet deploys yet - testnet only until the gates are reliable.
- I would rather show you a real beta with known issues than a polished demo that hides them.

Why the fit is 1:1

Injective runs	ArbiSim already runs
Native EVM (Chain ID 1776)	Anvil EVM fork engine - Injective support is a ChainConfig entry, same path as Avalanche and Arbitrum
ERC-8004 agent registry	ERC-8004 checks inside fork environments
x402 agent payments	x402 pay-per-simulation
MCP Server for agent trading	preflight_simulate as one more MCP tool call before the trade

What this gives Injective

A default safety layer under agent trading: fewer failed transactions, fewer burned users, and a claim no other chain can make - agent trades on Injective are safety-checked before execution. Two things I want to build here specifically:

- **Verdicts tied to ERC-8004 identities** - an agent's pre-flight history becomes part of its on-chain reputation, so the Injective registry gets richer with every simulation.
- **Public safety feed** - anonymized, aggregated risk data (which contracts flag, where sandwiches cluster) that the ecosystem can point to.

Who I am and what I am asking

I am a product manager, not a hardcore developer. I built this beta solo and I have led technical teams before - that is the part I am good at. What I am asking for is simple: if you think the idea is solid, help me build it on Injective. Tell me where it is wrong, point me at the right modules, and if it earns ecosystem support, that funds the developers to ship it properly. If the idea does not hold up under your questions, I want to know that too.

Links

Platform: arbisinguard.com | Docs: arbisinguard.com/docs | Analytics: dune.com/arbisinguard | X: [@Anteratic](https://twitter.com/Anteratic)